

## Depth Anything V2 Variants

ViT-S  
24.8M params

ViT-B  
97.5M params

ViT-L  
335.3M params

## Conversion Pipeline

PyTorch

ONNX  
(opset 17)

Simplified  
ONNX

TensorRT  
Engine

FP32

FP16

## Benchmark on NVIDIA RTX 4070 Laptop GPU

Latency  
(CUDA Events)

- mean, median
- p95, p99
- 200 iterations

Memory & Power  
(NVML)

- VRAM peak
- Avg power (W)

Accuracy  
(NYU Depth V2)

- AbsRel, RMSE
- $\delta_1$ ,  $\delta_2$ ,  $\delta_3$
- 654 samples